
You are currently looking at **version 1.0** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](https://www.coursera.org/learn/python-text-mining/resources/d9pwm) (<https://www.coursera.org/learn/python-text-mining/resources/d9pwm>), course resource.

Assignment 2 - Introduction to NLTK

In part 1 of this assignment you will use nltk to explore the Herman Melville novel Moby Dick. Then in part 2 you will create a spelling recommender function that uses nltk to find words similar to the misspelling.

Part 1 - Analyzing Moby Dick

```
In [1]: import nltk
import pandas as pd
import numpy as np

# If you would like to work with the raw text you can use 'moby_raw'
with open('moby.txt', 'r') as f:
    moby_raw = f.read()

# If you would like to work with the novel in nltk.Text format you can use 'text1'
moby_tokens = nltk.word_tokenize(moby_raw)
text1 = nltk.Text(moby_tokens)
```

Example 1

How many tokens (words and punctuation symbols) are in text1?

This function should return an integer.

```
In [2]: def example_one():

    return len(nltk.word_tokenize(moby_raw)) # or alternatively len(text1)

example_one()
```

Out[2]: 254989

Example 2

How many unique tokens (unique words and punctuation) does text1 have?

This function should return an integer.

```
In [3]: def example_two():

    return len(set(nltk.word_tokenize(moby_raw))) # or alternatively len(set(text1))

example_two()
```

Out[3]: 20755

Example 3

After lemmatizing the verbs, how many unique tokens does text1 have?

This function should return an integer.

```
In [4]: from nltk.stem import WordNetLemmatizer

def example_three():

    lemmatizer = WordNetLemmatizer()
    lemmatized = [lemmatizer.lemmatize(w, 'v') for w in text1]

    return len(set(lemmatized))

example_three()
```

Out[4]: 16900

Question 1

What is the lexical diversity of the given text input? (i.e. ratio of unique tokens to the total number of tokens)

This function should return a float.

```
In [5]: def answer_one():

        return # Your answer here

answer_one()
```

Question 2

What percentage of tokens is 'whale' or 'Whale'?

This function should return a float.

```
In [6]: def answer_two():

        return # Your answer here

answer_two()
```

Question 3

What are the 20 most frequently occurring (unique) tokens in the text? What is their frequency?

This function should return a list of 20 tuples where each tuple is of the form (token, frequency). The list should be sorted in descending order of frequency.

```
In [7]: def answer_three():

        return # Your answer here

answer_three()
```

Question 4

What tokens have a length of greater than 5 and frequency of more than 150?

This function should return a sorted list of the tokens that match the above constraints. To sort your list, use sorted()

```
In [8]: def answer_four():

        return # Your answer here

answer_four()
```

Question 5

Find the longest word in text1 and that word's length.

This function should return a tuple (longest_word, length).

```
In [9]: def answer_five():  
  
        return # Your answer here  
  
answer_five()
```

Question 6

What unique words have a frequency of more than 2000? What is their frequency?

"Hint: you may want to use `isalpha()` to check if the token is a word and not punctuation."

This function should return a list of tuples of the form (frequency, word) sorted in descending order of frequency.

```
In [10]: def answer_six():  
  
         return # Your answer here  
  
answer_six()
```

Question 7

What is the average number of tokens per sentence?

This function should return a float.

```
In [11]: def answer_seven():  
  
         return # Your answer here  
  
answer_seven()
```

Question 8

What are the 5 most frequent parts of speech in this text? What is their frequency?

This function should return a list of tuples of the form (part_of_speech, frequency) sorted in descending order of frequency.

```
In [12]: def answer_eight():  
  
         return # Your answer here  
  
answer_eight()
```

Part 2 - Spelling Recommender

For this part of the assignment you will create three different spelling recommenders, that each take a list of misspelled words and recommends a correctly spelled word for every word in the list.

For every misspelled word, the recommender should find the word in `correct_spellings` that has the shortest distance*, and starts with the same letter as the misspelled word, and return that word as a recommendation.

*Each of the three different recommenders will use a different distance measure (outlined below).

Each of the recommenders should provide recommendations for the three default words provided: ['cormulent', 'incendenece', 'validate'].

```
In [13]: from nltk.corpus import words
correct_spellings = words.words()
```

Question 9

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance (https://en.wikipedia.org/wiki/Jaccard_index) on the trigrams of the two words.

This function should return a list of length three: ['cormulent_reccommendation', 'incendenece_reccommendation', 'validate_reccommendation'].

```
In [14]: def answer_nine(entries=['cormulent', 'incendenece', 'validate']):

    return # Your answer here

answer_nine()
```

Question 10

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance (https://en.wikipedia.org/wiki/Jaccard_index) on the 4-grams of the two words.

This function should return a list of length three: ['cormulent_reccommendation', 'incendenece_reccommendation', 'validate_reccommendation'].

```
In [15]: def answer_ten(entries=['cormulent', 'incendenece', 'validate']):

    return # Your answer here

answer_ten()
```

Question 11

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Edit distance on the two words with transpositions.
(https://en.wikipedia.org/wiki/Damerau%E2%80%93Levenshtein_distance)

This function should return a list of length three: ['cormulent_reccommendation', 'incendenece_reccommendation', 'validate_reccommendation'].

```
In [16]: def answer_eleven(entries=['cormulent', 'incendenece', 'validate']):

    return # Your answer here

answer_eleven()
```